

# Erlang Concurrency and Distribution

Lars-Åke Fredlund  
2022



**POLITÉCNICA**

Master in Metodos Formales

Universidad Politécnica de Madrid

# Concurrency and Distribution?

According to Wikipedia:

- **Concurrent computing** is a form of computing in which several computations are executed concurrently – during overlapping time periods – instead of sequentially, with one completing before the next starts.

# Concurrency and Distribution?

According to Wikipedia:

- **Concurrent computing** is a form of computing in which several computations are executed concurrently – during overlapping time periods – instead of sequentially, with one completing before the next starts.
  - ▶ *Multiple activities happening at the same time*

# Concurrency and Distribution?

According to Wikipedia:

- **Concurrent computing** is a form of computing in which several computations are executed concurrently – during overlapping time periods – instead of sequentially, with one completing before the next starts.
  - ▶ *Multiple activities happening at the same time*
- A **distributed system** is a system whose components are located on different networked computers, which communicate and coordinate their actions by passing messages to one another from any system.

# Concurrency and Distribution?

According to Wikipedia:

- **Concurrent computing** is a form of computing in which several computations are executed concurrently – during overlapping time periods – instead of sequentially, with one completing before the next starts.
  - ▶ *Multiple activities happening at the same time*
- A **distributed system** is a system whose components are located on different networked computers, which communicate and coordinate their actions by passing messages to one another from any system.
  - ▶ *Multiple activities happening at different places*

# Concurrency and Distribution in Erlang: Highlights

- ❶ An **intuitive programming model** for concurrency – matching some common sense intuitions we have as humans
- ❷ Using **concurrency is efficient**: concurrent programs can use available processing power well (e.g. making efficient use of multiple processor cores in a CPU)
- ❸ **Minimizing the differences between distributed and concurrent programming** – the same programming constructs are used for both

## An intuitive programming model: **Actors** (Carl Hewitt)

- A concurrent system is composed of *Actors*
- An actor is a computation entity that can:
  - ▶ create new actors
  - ▶ send messages to other actors
  - ▶ act upon messages received from other actors, when the actor itself chooses to do so
- Compare with normal human behaviour...

## An intuitive programming model: **Actors** (Carl Hewitt)

- A concurrent system is composed of *Actors*
- An actor is a computation entity that can:
  - ▶ create new actors
  - ▶ send messages to other actors
  - ▶ act upon messages received from other actors, when the actor itself chooses to do so
- Compare with normal human behaviour...
- And compare with Java



## An intuitive programming model: **Actors** (Carl Hewitt)

- A concurrent system is composed of *Actors*
- An actor is a computation entity that can:
  - ▶ create new actors
  - ▶ send messages to other actors
  - ▶ act upon messages received from other actors, when the actor itself chooses to do so
- Compare with normal human behaviour...
- And compare with Java

The basic ideas of Actors have been implemented in different programming languages and libraries, for instance:

- *Akka* for Java and Scala
- *CAF* for C++
- *Thesbian* for Python
- ...

# Actors in Erlang

An actor is a....

An actor is a.... **process**

- A process has a unique *name* – its **process identifier** or **pid**
- A process executes a function call  $f(e_1, \dots, e_n)$ . When the function call terminates, the process terminates.
- A process has a **mailbox** where messages sent to the process are stored. The mailbox is a queue – messages received earlier are placed before messages received later.

Note that the Erlang shell `erl` executes all code inside a process.

## Interacting with processes in Erlang

- **What is my pid** (process identifier)?

```
self()
```

- **Create a new process** executing 1+2

```
spawn(fun () -> 1+2 end)
```

- **Sending** a message {hello, world} to a process with pid Pid:

```
Pid ! {hello, world}
```

- ▶ sending is *asynchronous, non-blocking*

- **Retrieving a message** from my mailbox:

```
receive  
  X -> X  
end
```

## The receive statement

```
receive  
   $pat_1$  when  $guard_1 \rightarrow expr_1$   
  ...  
   $pat_n$  when  $guard_n \rightarrow expr_n$   
  after  $time \rightarrow expr$   
end
```

- The **receive** statement removes a message from the mailbox
- A removed message has to:
  - ▶ match a clause pattern,
  - ▶ with the clause guard true,
  - ▶ *and no older message can be removed*
- **receive** returns the value computed by the expression in the **first** matching clause
- If no message can be removed the receive statement waits (potentially forever) for the reception of a removable message by the process
- The optional timeout clause labelled with **after** specifies a wait timeout

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  Other -> wrong  
end
```

and the queue is  $\{\text{inc}, 3\} \cdot \text{"a"}$  where  $\{\text{inc}, 3\}$  is the oldest message – what happens?

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  Other -> wrong  
end
```

and the queue is  $\{\text{inc}, 3\} \cdot \text{"a"}$  where  $\{\text{inc}, 3\}$  is the oldest message – what happens?

- 4 is returned
- The resulting queue is "a"

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  Other -> wrong  
end
```

and the queue is "b" · {inc, 3} · "a" where "b" is the oldest message – what happens?



## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  Other -> wrong  
end
```

and the queue is "b" · {inc, 3} · "a" where "b" is the oldest message – what happens?

- wrong is returned
- The resulting queue is {inc, 3} · "a"

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
end
```

and the queue is "b" · {inc, 3} · "a" where "b" is the oldest message – what happens?

## Receive – examples

Given a receive statement

```
receive  
    {inc, X} -> X+1;  
end
```

and the queue is "b" · {inc, 3} · "a" where "b" is the oldest message – what happens?

- 4 is returned
- The resulting queue is "b" · "a"

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
end
```

and the queue is "b" · "a" where "b" is the oldest message – what happens?

## Receive – examples

Given a receive statement

```
receive  
    {inc, X} -> X+1;  
end
```

and the queue is "b" · "a" where "b" is the oldest message – what happens?

- The receive statement waits until a matching message is received (potentially forever)

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  after 1000 -> {wrong, timeout}  
end
```

and the queue is "b" · "a" where "b" is the oldest message – what happens?

## Receive – examples

Given a receive statement

```
receive  
  {inc, X} -> X+1;  
  after 1000 -> {wrong, timeout}  
end
```

and the queue is "b" · "a" where "b" is the oldest message – what happens?

- The receive statement waits at most 1000 milliseconds until a matching message is received, and if no message is received returns {wrong, timeout}

## Exercise

How can we implement a function `sleep(time)` which sleeps for `time` milliseconds?



## Exercise

How can we implement a function `sleep(time)` which sleeps for `time` milliseconds?

```
sleep(time) ->  
  receive  
    after Time -> ok    % or some other return value...  
  end  
end.
```

## Exercise

- Read either the message "a" or the message "b" from the mailbox.
- If there is a message "a" always read that first.
- Only if there is no message "a" the message "b" may be read.
- If neither "a" nor "b" can be read wait until they can.

## Exercise

- Read either the message "a" or the message "b" from the mailbox.
- If there is a message "a" always read that first.
- Only if there is no message "a" the message "b" may be read.
- If neither "a" nor "b" can be read wait until they can.

```
receive
  a -> ...
after 0 ->
  receive
    a -> ...
    b -> ...
  end
end
```

## Concurrency – Extra Features: Naming a Process

- We can register a symbolic name for a process using **register**(Pid,Name)
- The symbolic name can be used when sending messages:

```
spawn(fun () ->
    register(server,self()),
    receive
        Msg -> io:format("Hola ~p",[Msg])
    end
end).

server ! "pablo".
```

- If a process  $P$  first sends the message  $m_1$  to  $Q$ , and then sends the message  $m_2$  to  $Q$ , what can we as programmers assume?

- If a process  $P$  first sends the message  $m_1$  to  $Q$ , and then sends the message  $m_2$  to  $Q$ , what can we as programmers assume?
  - ▶ can the messages be dropped (i.e., never arrive at  $Q$ )?

- If a process  $P$  first sends the message  $m_1$  to  $Q$ , and then sends the message  $m_2$  to  $Q$ , what can we as programmers assume?
  - ▶ can the messages be dropped (i.e., never arrive at  $Q$ )?
  - ▶ can the messages be duplicated (i.e.,  $m_1$  arrives twice at  $Q$ )?

- If a process  $P$  first sends the message  $m_1$  to  $Q$ , and then sends the message  $m_2$  to  $Q$ , what can we as programmers assume?
  - ▶ can the messages be dropped (i.e., never arrive at  $Q$ )?
  - ▶ can the messages be duplicated (i.e.,  $m_1$  arrives twice at  $Q$ )?
  - ▶ can they be corrupted? (i.e.,  $Q$  receives  $m_1'$  instead of  $m_1$ )




- If a process  $P$  first sends the message  $m_1$  to  $Q$ , and then sends the message  $m_2$  to  $Q$ , what can we as programmers assume?
  - ▶ can the messages be dropped (i.e., never arrive at  $Q$ )?
  - ▶ can the messages be duplicated (i.e.,  $m_1$  arrives twice at  $Q$ )?
  - ▶ can they be corrupted? (i.e.,  $Q$  receives  $m_1'$  instead of  $m_1$ )
  - ▶ what is the order in which  $Q$  receives these messages?
- Guarantee: **messages sent from a process  $P$  to a process  $Q$  are delivered in order (or  $P$  or  $Q$  crashes)**
- Concretely: if  $P$  sends first  $m_1$  and then  $m_2$  to  $Q$ , then first  $m_1$  will be stored in the mailbox of  $Q$ , and then  $m_2$  will be stored in its mailbox.

## Communication Guarantees II

Consider the following situation:

- 1  $P$  sends a message  $m_1$  to  $Q$
- 2  $P$  next sends a message  $m_2$  to  $R$
- 3  $R$  forwards the message  $m_2$  to  $Q$




In which order will  $Q$  receive the messages  $m_1$  and  $m_2$ ?

## Communication Guarantees II

Consider the following situation:

- 1  $P$  sends a message  $m_1$  to  $Q$
- 2  $P$  next sends a message  $m_2$  to  $R$
- 3  $R$  forwards the message  $m_2$  to  $Q$



In which order will  $Q$  receive the messages  $m_1$  and  $m_2$ ?

- We don't know - either  $m_1$  or  $m_2$  will arrive first in the mailbox of  $Q$
- Mimics TCP/IP communication guarantees

# Handling Errors – Process Termination

A process can die for different reasons:

- It can terminate *normally* – e.g.,

```
spawn(fun () -> fib(10) end)
```

- It can terminate *abnormally* because of an exception which was not handled:

- ▶ Common mistakes: `2/0`
- ▶ Signalling errors: `error("do not do that")`
- ▶ Non-caught throws: `throw(hola)`

- Or it can terminate *abnormally* if the process was ordered to die:

```
exit(Pid, Reason)
```

- Suppose  $P$  spawns a new process  $Q$ , and  $Q$  dies *abnormally*, what happens to  $P$ ?

## Handling Errors: Links

- Suppose  $P$  spawns a new process  $Q$ , and  $Q$  dies *abnormally*, what happens to  $P$ ?
  - ▶ By default absolutely nothing

## Handling Errors: Links

- Suppose  $P$  spawns a new process  $Q$ , and  $Q$  dies *abnormally*, what happens to  $P$ ?
  - ▶ By default absolutely nothing

- Processes can be *linked* together using a call **link**(Pid):

```
Pid = spawn(fun () -> ... end),  
link(Pid).
```

- The process executing the spawn call will be *linked* to the new spawned process

## Handling Errors: Links

- Suppose  $P$  spawns a new process  $Q$ , and  $Q$  dies *abnormally*, what happens to  $P$ ?
  - ▶ By default absolutely nothing

- Processes can be *linked* together using a call **link**(Pid):

```
Pid = spawn(fun () -> ... end),  
link(Pid).
```

- The process executing the spawn call will be *linked* to the new spawned process
- We can spawn and link at the same time – **atomically** – using

```
Pid = spawn_link(fun () -> ... end)
```



## Handling Errors: Links and Abnormal Termination

Suppose we execute the following in process  $P$ :

```
Pid = spawn_link(fun () -> ... end)
```

- What happens if the spawned process dies *abnormally*?

## Handling Errors: Links and Abnormal Termination

Suppose we execute the following in process  $P$ :

```
Pid = spawn_link(fun () -> ... end)
```

- What happens if the spawned process dies *abnormally*?
  - ▶ **By default  $P$  dies *abnormally* too!**

## Handling Errors: Links and Abnormal Termination

Suppose we execute the following in process  $P$ :

```
Pid = spawn_link(fun () -> ... end)
```

- What happens if the spawned process dies *abnormally*?
  - ▶ **By default  $P$  dies *abnormally* too!**
- What happens if the spawned process dies *normally*?

## Handling Errors: Links and Abnormal Termination

Suppose we execute the following in process  $P$ :

```
Pid = spawn_link(fun () -> ... end)
```

- What happens if the spawned process dies *abnormally*?
  - ▶ **By default  $P$  dies *abnormally* too!**
- What happens if the spawned process dies *normally*?
  - ▶ By default nothing happens to  $P$ .

Suppose we execute the following in process  $P$ :

```
Pid = spawn_link(fun () -> ... end)
```

- What happens if the spawned process dies *abnormally*?
  - ▶ **By default  $P$  dies *abnormally* too!**
- What happens if the spawned process dies *normally*?
  - ▶ By default nothing happens to  $P$ .
- Note that links are *bidirectional*. Suppose  $P$  and  $Q$  are linked:
  - ▶ if  $P$  dies abnormally so does  $Q$
  - ▶ if  $Q$  dies abnormally so does  $P$ !

## Handling Errors: Links and Abnormal Termination as Messages

- As an option the spawned process can be *notified* instead of terminated when a linked process terminates

## Handling Errors: Links and Abnormal Termination as Messages

- As an option the spawned process can be *notified* instead of terminated when a linked process terminates
- Notifications are received as normal messages in the mailbox on the format

`{'EXIT', Pid, Reason}`

where `pid` is the pid that terminated and `reason` is the reason for termination

- The notification behaviour is chosen by calling **`process_flag(trap_exit, true)`**:

```
process_flag(trap_exit, true),  
Pid = spawn_link(fun () -> ... end).
```

## Concurrency – Extra Features: Monitors

- *Monitors* behave like links, **except** they:
  - ▶ are unidirectional
  - ▶ **always** notifies (sends termination messages) – **never** propagates abnormal termination

```
Pid = spawn(fun () -> 2 end),  
monitor(Pid)
```



# Concurrency in Erlang is Efficient

- Pros:

- ▶ It is cheap (quick, uses little memory) to spawn processes
- ▶ It is quick to switch which process is executing on a processor core
- ▶ The concurrency model maps well onto today's hardware: makes use of multiple processor cores with *no programming overhead* and *little computing overhead*

# Concurrency in Erlang is Efficient

- Pros:

- ▶ It is cheap (quick, uses little memory) to spawn processes
- ▶ It is quick to switch which process is executing on a processor core
- ▶ The concurrency model maps well onto today's hardware: makes use of multiple processor cores with *no programming overhead* and *little computing overhead*

- Cons:

- ▶ Sending large messages between processes is not cheap (copying is necessary)
- ▶ Thus data structures shared between multiple processes are sometimes used (**the horror!**).  
Example: the ets term storage library.

# Concurrency in Erlang is Efficient

- Pros:

- ▶ It is cheap (quick, uses little memory) to spawn processes
- ▶ It is quick to switch which process is executing on a processor core
- ▶ The concurrency model maps well onto today's hardware: makes use of multiple processor cores with *no programming overhead* and *little computing overhead*

- Cons:

- ▶ Sending large messages between processes is not cheap (copying is necessary)
- ▶ Thus data structures shared between multiple processes are sometimes used (**the horror!**).  
Example: the ets term storage library.

- In practice is realistic to have systems with up to 100,000 alive processes

- Erlang “nodes” – Erlang runtime systems – can be connected. API:
  - ▶ Provide a name for the erl runtime at startup: “erl -sname nodename”
  - ▶ **node()** – returns the node name
  - ▶ **net\_adm:ping(nodename)** – connect to the remote node nodename
- New processes can be spawned at a connected remote node:  
Node:**spawn**(nodename, **fun** ()->... **end**)
- **Process-to-process communication, links, etc, work between processes at different nodes with no syntactic differences!**

# Support for Distribution in Erlang: Pros and Cons

- Positive
  - ▶ a very simple and powerful programming model

# Support for Distribution in Erlang: Pros and Cons

- Positive
  - ▶ a very simple and powerful programming model
- Negative
  - ▶ communication between nodes is not encrypted
  - ▶ performance is realistic for only small networks of nodes. The network is *fully connected*: every node has an explicit connection to every other node in the network
  - ▶ difficulty of integrating non-Erlang nodes in the network

# Support for Distribution in Erlang: Pros and Cons

- Positive
  - ▶ a very simple and powerful programming model
- Negative
  - ▶ communication between nodes is not encrypted
  - ▶ performance is realistic for only small networks of nodes. The network is *fully connected*: every node has an explicit connection to every other node in the network
  - ▶ difficulty of integrating non-Erlang nodes in the network
- In practice
  - ▶ used only for **small internal networks**, e.g., to distribute load between a group of servers
  - ▶ often non-Erlang “middlewares” are used to enable communication between Erlang programmed nodes (RabbitMQ, ...)